

Measuring what an LLM adds to malware-to-ATT&CK technique mapping

Düzgün Küçük^{a,*}, Fatih Ertam^a

^a*Department of Digital Forensics Engineering, Faculty of Technology, Fırat University, Elazığ, Türkiye*

Abstract

A large language model (LLM) pipeline for mapping malware evidence to MITRE ATT&CK (Adversarial Tactics, Techniques and Common Knowledge) techniques is mostly not the model. It wraps one in a sandbox, a disassembler, retrieval indices and rule engines. What the model contributes is therefore a difference against that tooling, not a score. This study measures that difference component by component. Three findings are positive. Splitting the analysis across several model calls scores +0.0537 F1 above a single judge on 24 real samples [+0.0362, +0.0714], at 2.76 times its output. A deterministic technique assigner that composes its ranking and gating backends, rather than choosing between them, gives the cleanest identifier gate on two external corpora. And turning one decoding flag off gains 0.34 to 0.45 F1, more than any architecture or parameter count we measured. Four findings are negative or bounded. The negotiation the design was built for contributes +0.0005, at a resolution of 0.046: what pays is the extra calls, not the argument between them. Against the sandbox it is built on, the full pipeline is +0.0030 F1 [-0.0353, +0.0344], at a resolution of 0.085. Three retrieval components move nothing end to end. And the confidence the model reports, which every deterministic gate consumes, separates correct claims from wrong ones at an area under the curve of 0.550, on an interval containing chance. The paper also reports seven defects in its own instrument, each of which returned a plausible result rather than an error.

Keywords: LLM agents, multi-agent systems, malware analysis, MITRE ATT&CK, evaluation methodology, silent failure, negative results

1. Introduction

Mapping malware evidence to the technique identifiers of MITRE ATT&CK (Adversarial Tactics, Techniques and Common Knowledge) is a labelling task that large language models (LLMs) are now routinely applied to, and the systems that

*Corresponding author.

Email address: dzgnkucuk@gmail.com (Düzgün Küçük)

perform it are rarely a model alone. A working pipeline wraps the model in deterministic tooling: a sandbox that detonates the sample, a disassembler that recovers its structure, retrieval indices that supply prior cases, and rule engines that assert techniques directly from signatures. The model’s contribution is therefore a difference against that tooling rather than a score in isolation, and that difference is almost never reported. Büchel et al. [1] re-evaluate more than forty technique-extraction systems in one setting and find them “largely incomparable”, because each uses its own ontology and its own inaccessible dataset. An F1 value read against no baseline refers to nothing, and the baseline that matters for an LLM pipeline is the deterministic engine it was built on top of.

A second barrier is that such a pipeline is itself a distributed system, and its evaluation inherits every property of one. Each boundary between the pipeline and a server it depends on is a place where a request can succeed and still not mean what the caller assumed, and the resulting error is not visible as an error. Silent failure of this kind is a long-established class in production software [2, 3, 4] and is now being catalogued for LLM orchestration frameworks and inference engines as well [5, 6]. What changes in an evaluation setting is the artefact. The product of a silent failure is not a degraded user session that someone reports; it is a measurement that is wrong, looks right, and is written down.

This work reports the evaluation of one such pipeline, built for ATT&CK mapping over static, dynamic and network evidence, and the failures of the instrument that produced that evaluation. Figure 1 shows the system and Section 2.2 describes it. It was built around four architectural claims: that multi-agent consensus over heterogeneous evidence channels beats a single reader, that a multi-layer corroboration cascade improves evidence quality, that two-tier attribution supplies useful priors, and that falsification before confidence disciplines unsupported claims. Each claim was measured under the design of Section 2: an equal token budget against a simpler alternative, a stochastic-noise control, and intervals resampled at the level the observations are independent at.

One of those four survives in part and three do not. Section 1.2 states each finding with the effect the design that tested it could have detected.

The study that established all of this was itself wrong several times over. Seven defects in the instrument produced plausible results rather than errors, and 2,716 passing tests caught none of them.

1.1. Related work

Language models are applied across the malware task surface: tracking malicious packages in a registry [7], detecting Android malware from application features [8] and from generated training samples [9], and summarising a sample’s behaviour in prose [10]. None produces ATT&CK technique identifiers from a running analysis of a Windows binary. Technique extraction itself has moved from supervised classifiers, among them the Threat Report ATT&CK Mapper (TRAM) [11], rcATT [12] and TTPDrill [13], through neural extractors such as EXTRACTOR [14], AttackKG [15] and LADDER [16], to retrieval-augmented generation

(RAG). TechniqueRAG [17] fine-tunes only the generator, because ATT&CK defines over 550 sub-techniques against roughly 10,000 public annotated examples. A hierarchical successor cuts the candidate space by 77.5% [18] [preprint], and TTPrint [19] [preprint] keeps only candidates supported by both an evidence span and the MITRE definition. Neither is a comparator here: both are unreviewed, and both report against baselines that Büchel et al. [1] call “largely incomparable” in a survey re-evaluating more than forty systems, which also finds traditional natural language processing (NLP) approaches outperforming embedder-based and generative ones. Our describe-then-map design removes technique-identifier recall from the model entirely, a domain instantiation of Infer-Retrieve-Rank [20] [preprint] stricter than TechniqueRAG [17], where the model still ranks among retrieved candidates. That a retrieval backend can rank well and gate badly is likewise not ours: Abdallah et al. [21] treat retrieval confidence as distinct from effectiveness and find calibration “consistently weak across model families”. Section 3.2 measures that separation on TRAM2 [11] and on the independently annotated AnnoCTR corpus [22]. Two further results we rely on have prior art we cite rather than compete with: that a false correction costs more than a missed one, which is why grammatical error correction evaluates with F0.5 [23, 24], and that neural text degeneration has two competing accounts, one in the likelihood objective and one in training-data repetition [25, 26].

Agentic systems drive reverse-engineering backends through tool interfaces, and the models split by role: the binary specialists are small, open-weight and fine-tuned, among them ReCopilot [27] [preprint], LLM4Decompile [28] and Nova [29], while the models that drive the tools are frontier models. TTPDetect [30] [preprint] is the closest system to ours, identifying techniques in stripped binaries through retrieval-narrowed entry points and a function-level agent at 93.25% precision and 93.81% recall at function level and 87.37% precision on real-world malware; we do not claim binary evidence as an input modality, and that ground is theirs. Tool selection degrades as a catalogue grows [31] [preprint] and narrowing one recovers execution time [32], though neither result covers tools that overlap in what they do. Running the model locally for confidentiality is established: REx86 [33] fine-tunes open-weight models for x86 reverse engineering because “cloud-hosted, closed-weight models pose privacy and security risks and cannot be used in closed-network facilities”.

Multi-agent debate is a standard pattern, including in security [34]. Two recent results change how such claims must be read. Holding the reasoning budget constant, Tran and Kiela find single agents match or beat multi-agent systems on multi-hop reasoning, with a crossover appearing only at degradation 0.7 [35] [preprint], and Bertalanič and Fortuna [36] reach the same conclusion on 7B and 8B open-weight models, where debate consumed 2.1 to 3.4× more tokens for equal or lower accuracy. Both scope their negatives to homogeneous agents splitting one context, and both name degraded or noisy context as the exception. That exception is our setting: a 600-import binary, a sandbox report and a network capture are three channels rather than one context divided, and our equal-budget ablation asks whether

splitting by channel survives a control that splitting by context does not. Weighting evidence by how far its source is trusted, and discounting sources that are not independent, is Dempster-Shafer theory; our cascade is an instance of it with hand-chosen constants, and Section 3.6 measures how far those constants reach.

This work’s methodological results belong to an established critique tradition. Arp et al.’s *Dos and Don’ts* identified ten pitfalls in machine learning for security, and its successor *Chasing Shadows* [37] surveys all 72 peer-reviewed LLM-security papers of 2023 and 2024, finding every one exposed to at least one of nine pitfalls and only 15.7% discussing any. A negative RAG result in malware analysis is already published by a different mechanism, retrieved context diluting a signal-extraction task [38]. Comparing against a trivial label-only baseline is long established [39], and a systematic review by the same group found a considerable share of published multi-label results no better than it [40] [preprint]. On temporal stability the axis matters more than the result: performance falls as evaluation moves away from the training period and scale does not close the gap [41]. That work varies the gap between training and test time; ours fixed the model and varied the era of the input binary, and we no longer have a result on that axis (Section 3.11).

The paper’s methodological spine belongs to an area that already exists. Gunawi et al. [2] measured where an error signal is lost, finding 13% of propagation channels drop an error code without handling it; Yuan et al. [3] attributed 25% of catastrophic failures to errors explicitly signalled and then ignored; and Alquraan et al. [4] found the majority of network-partition failures in cloud systems to be silent. The LLM-specific instance is now catalogued in bug taxonomies for agent orchestration frameworks [5] and inference engines [6], the second naming both incorrect request parsing as a cause and silent error as a symptom. Overruling a requested parameter back to a default without saying so is itself a named misconfiguration failure mode [42]. Mining Model Context Protocol (MCP) reference servers, Oyelajo et al. [43] find data validation and type issues to be the second most prevalent defect class at 13.97%, which is the class our own first defect falls into; a related constraint is that an application programming interface (API) specification cannot express every dependency between its inputs, such dependencies being “the norm, rather than the exception, with 85% of the reviewed APIs” [44]. The detector is not new either: “distinct inputs should produce distinct outputs” is an ordinary metamorphic relation [45], and stress-testing an LLM judge’s consistency under input variation is established [46]. Nor is the newest of our mechanisms: our own parameter-size series reproduced the $\rho=0.85$ at $p=0.014$ the nearest prior work reports [47] [preprint], on an axis whose low-scoring row was a model disabled by a configuration flag rather than limited by its size. That a setting can outweigh model size is established, under output-length constraints [48], across inference-time configurations [49], and where extra test-time compute lets a smaller model overtake a larger one [50, 51] or longer reasoning improves accuracy [52].

Every result of ours cited above was produced on one model on one machine, except where a statement concerns retrieval or the deterministic cascade and no model was involved at all.

1.2. Motivation and contributions

A component of an LLM analysis pipeline is ordinarily justified by the mechanism it implements rather than by a measurement, and a mechanism that is sound in isolation can contribute nothing once wired to real inputs. Establishing which of the two holds needs an equal-budget comparison against a simpler alternative, and a baseline containing no language model at all. The contributions are therefore organised around the evaluation rather than the architecture, and each is a measurement rather than a proposal.

(i) *A deterministic technique assigner, measured on two external corpora.* Ranking quality and gate quality are separable axes on which no single backend wins, and composing the per-axis winners gives the cleanest identifier gate of the three on both TRAM2 and AnnoCTR, with all three orderings replicating across corpora annotated by different people for different tasks (Section 3.2).

(ii) *What decomposition is worth, and what it is not worth.* On 24 real samples, decomposing the analysis into several model calls beats a single judge by +0.0537 F1 at 2.76 times its output, while the contrast that isolates the negotiation returns +0.0005 at matched calls: what pays is the calls, not the argument between them.

(iii) *A no-LLM baseline for ATT&CK mapping, and the pipeline scored against it.* Scored per sample against family-level MITRE ground truth, the signature engine beneath the pipeline reaches F1 0.1526 [+0.1114, +0.1998] on 97 samples, and on the samples where both arms ran the full pipeline is +0.0030 F1 above it at a resolution of 0.085.

(iv) *Three architectural claims that did not survive, each with the resolution of the design that tested it.* A weighted corroboration cascade, two-tier attribution and a deterministic confidence gate each returned a negative or near-null result, stated with the effect its design could have detected (Section 2.6). The configuration axis is where the effects are: a decoding flag outweighs every architecture and every parameter count we measured, and an arm can be labelled matched while running the opposite configuration.

(v) *Seven instrument failures, with the layer each occurred at and what actually found it.* Each is a case where a measurement's cause was assigned rather than measured.

2. Materials and methods

2.1. Scope and assumptions

This study evaluates one pipeline that maps evidence about a Windows Portable Executable (PE) malware sample to MITRE ATT&CK technique identifiers and emits a Structured Threat Information Expression (STIX) bundle. The scope is narrow in three respects and every result should be read inside it. The model is a single local mixture-of-experts deployment, quantised, on one host, with three commercially hosted services used only as comparison endpoints. Ground truth is family-level MITRE uses, resolved to a sample through its family signature, which

is coarse and carries a structural recall ceiling. Only Windows PE is analysed dynamically, because the sandbox instance has one analysis machine and no Linux guest.

2.2. The pipeline under evaluation

Figure 1 shows the system. Its organising rule is that the model proposes and a deterministic layer disposes: the model describes behaviour and never emits a technique identifier or a final set, and every identifier that reaches the analyst has been assigned, gated and reconciled by code. That is why the deterministic layer dominating the output is the design working rather than a component failing, and why what has to be measured is what the model adds on top of it. Six deterministic evidence layers assert techniques directly from signatures, rules and byte markers, with no model involved. Three analysts then run as ReAct loops, each bound to one tool server and reading one channel of evidence. A negotiation node tests for consensus and routes disputes to a revision pass. A judge then synthesises a verdict and emits STIX 2.1. Finally, a deterministic reconciliation and gating stage adjudicates between the judge’s bundle and the corroboration cascade, before the analyst sees anything. The prompts and the block format an analyst answers in are in the released artefact [53]. The rule the three share is the pipeline’s grounding constraint: every claim must cite a concrete artefact, a function name, an API call or a resolved domain, and a claim may carry a confidence above 0.8 only after a falsification check the model is told to run. That is the discipline whose output Section 3.5 finds near chance.

The cascade scores a technique t by weighting each contributing layer’s confidence by how far that layer is trusted, then lifting the result by how many layers agree:

$$s(t) = \min\left(\frac{\sum_{d \in D(t)} w_d c_d(t)}{\sum_{d \in D(t)} w_d} m(|D(t)|), 1\right) \quad (1)$$

$D(t)$ are the domains claiming t and $c_d(t)$ is domain d ’s confidence for it, a mean over a model’s claims and a maximum over a rule’s. The trust weights w run from 0.90 for YARA to 0.20 for network evidence, and the corroboration multiplier m rises from 1.00 at one contributing layer to 1.90 at 5.

Two design choices are not what failed. Tool exposure is not tool availability: exposing a full reverse-engineering catalogue to a model with roughly 3B active parameters is infeasible, because tool selection degrades before the context does. The static analyst is therefore given a curated allowlist of 20 tools, and the high-value tools are invoked deterministically from code. And the pipeline degrades rather than fails when a service is absent, so with the sandbox unreachable the dynamic path is skipped and the run completes on static evidence, a behaviour pinned by a test. That mattered during the study, because the sandbox was reachable from one network only and every static-only result reported here was produced by that degradation working as designed.




Figure 1: The pipeline under evaluation, with each of the seven instrument failures of Section 3.8 marked at the boundary it occurred at. Trust weights and cross-layer multipliers are the shipped constants.

2.3. *Samples, ground truth and the sandbox*

The dynamic cohort is $n=100$ Windows PE samples from MalwareBazaar, stratified by year at seed 20260810 with its digest recorded, of which 97 produced a usable sandbox report. Ground truth resolves a family signature to an in-repo MITRE uses fixture through an alias map, and is uncurated at technique level, so a single binary of a family need not exhibit all the techniques catalogued for it: absolute F1 values here are not comparable to work using per-sample expert labels, and the bias being approximately constant across arms is what makes within-study contrasts the defensible reading. The withdrawn temporal-drift study’s manifest survives as 210 dated samples over 7 cohorts and 27 families, metadata only. Three properties of the sandbox bound what can be reproduced. It is CAPE 2.5, the Config and Payload Extraction sandbox, with one analysis machine and no Linux guest, so an Executable and Linkable Format binary submitted there joins a backlog that is never scheduled. Reports are retained for days rather than indefinitely, 18 of 18 sampled older tasks reporting no report directory, so the local archive is the only copy. And the instance rate-limits with a throttle response byte-identical to its refusal of an unauthorised request, so a harness reading that response as a verdict on its own permissions mistakes a throttle for a missing permission, as ours did briefly.

2.4. *Models, endpoints and serving configuration*

Table 1 identifies the local model and inference engine by Secure Hash Algorithm (SHA-256) digest and by source commit. Naming the imatrix calibration dataset matters, because two IQ3 quantisations of the same model calibrated differently are not the same artefact. The engine commit was reconstructed rather than recorded: the running binary reports no version, on a build tree that was never a git checkout, so the commit was recovered from a vendored clone and proved against an identical source list.

Table 2 lists the endpoints each arm ran against. One parameter on it decides more than the model names do. `enable_thinking` is honoured on DashScope and ignored on OpenRouter. Ignored is the exact word: the request is accepted, the value is recorded in the result file, and nothing acts on it, so the arm that asked for reasoning and the arm that did not spend the same share of their output on it (Section 3.4). The flag outweighs the differences the arms were built to compare, so an arm can be labelled matched while running the opposite configuration, and arm selection therefore keys on the measured reasoning share of a completed arm rather than on the flag requested.

Two serving details decide whether an arm reproduces, and the invocation itself is in the released artefact [53]. The local server’s output cap must be sent twice, because the client library renames it on the wire to a spelling `ik_llama.cpp` does not read (Section 3.8); and that server truncates silently, reporting the same finish reason at the cap as on a completed answer, so telemetry keyed to the finish reason reports no truncation however often the cap binds. The frontier arm ran on a rate-limited free tier whose exhaustion returned HTTP 429, the Hypertext Transfer

Table 1: The local model and inference engine, pinned by digest and commit, and the four measured host properties that bound what can be reproduced on it.

model	Qwen/Qwen3.6-35B-A3B, quantised IQ3_K_R4 by Unsloth
model file SHA-256	d0de70ef...c4ea, computed locally and matching the HuggingFace etag
HuggingFace revision	cf350fd...1f0d, retrieved 2026-05-11
imatrix calibration	unsloth_calibration_Qwen3.6-35B-A3B.txt
architecture	hybrid recurrent, from the model file metadata
engine	ik_llama.cpp at eb570eb9...d26de
host	one RTX 5060, 31 GiB random-access memory, 16 threads
server growth	10.4 GB fresh, 14.8 GB after one analyst pass, 18.1 GB after 60 minutes; anonymous and unreclaimable, so long paired runs restart the server between arms
context	131,072 of the native 262,144, 30 blocks on the central processing unit; serving 64k instead does not lower the peak (14.6 against 14.8 GB)
generation rate	162 to 20 tokens/s, because the engine writes and erases a 63 MiB recurrent-state checkpoint every few hundred tokens
disassembly container	capped at 6 GB, its Java virtual machine flag bounding the heap only, measured resident set size 5.15 GB against that nominal 4 GB cap

Table 2: The inference endpoints each arm ran against, with the configuration measured rather than the configuration requested.

arm	endpoint	model	reasoning	n
local	ik_llama.cpp, this host	Qwen3.6-35B-A3B (IQ3_K_R4)	off	25
baseline	OpenRouter	nvidia/nemotron-3-super-120b-a12b:free	on, not controllable	25 × 2
frontier	DashScope	qwen3.6-35b-a3b	off	25
same weights, hosted	DashScope	qwen3.6-35b-a3b	on	25
same weights, hosted	DashScope	qwen3.6-35b-a3b	on	25
larger, hosted	DashScope	qwen3.6-plus	off / on	25 each

Protocol (HTTP) status for a rate limit, which is why that arm is 25 calls rather than the 97-sample cohort.

2.5. Experimental design

Four terms are used in one sense throughout. An *arm* is one condition of one experiment, a configuration under which every input is scored, so two arms differ in the one thing being tested. A *sample* is a real malware binary from the dated cohort, scored against its family’s MITRE uses set. A *fixture* is one of the five synthesised inputs used by the equal-budget arms, its evidence generated from its own technique list, so it has a ceiling a real sample does not; samples and fixtures are different populations with different ground truth, never pooled and never drawn on one axis. The *reconciliation step* is the deterministic pass that compares the judge’s bundle against the corroboration cascade’s set and restores what the judge omitted.

Arms are capped at an equal total token budget: a multi-agent arm making N calls receives B/N per call against a single arm’s one call of B . That cap is a ceiling rather than an allocation, and the single arm does not spend its own, so the measured output is $3.2\times$ the single arm’s on the fixtures and $2.76\times$ on the real samples. The arms are matched on what they were allowed rather than on what they used. A contrast against the single arm therefore cannot separate the decomposition from the generation the extra calls bring, and only the contrast between the two multi-agent arms, which make the same number of calls, isolates a mechanism. The design follows Bertalaníč and Fortuna [36], including a stochastic-noise control in which one analyst receives evidence irrelevant to the sample. That control is what distinguishes no effect from no sensitivity: an arm fed evidence about the wrong sample should separate, and a design in which it does not cannot be read as having found nothing. For a reasoning model the budget must cover the reasoning as well. Our frontier endpoint spent 56.5% of its output on reasoning across 25 real prompts, and 84% on a one-token answer, so capping the answer alone would have granted it roughly twice the generation for the same nominal budget. Arms must also sit on the same side of the reasoning switch, because a model that spends its whole output budget thinking returns no answer at all: an unmatched pair does not produce a noisy comparison, it produces a comparison of the flag (Section 3.4).

Three further rules were adopted after each one caught something, and all three are cheap. The first is that a mechanism’s firing rate is measured before its effect, because averaging a feature over inputs it never touched dilutes a real effect toward zero: an ablation has to be restricted to the samples where the mechanism fires, and a mechanism that fires on almost nothing cannot be ablated at all (Section 3.5). The second is that the number of distinct outputs the N inputs produced is counted and reported beside the sample size rather than asserted in a test (Section 3.9). The third is that per-sample outputs are retained for every study, since aggregates cannot be re-interrogated: when a defect was found that might have affected a completed 210-sample study, the question could not be asked because only the summary survived. Long-lived servers accumulate state that crosses sample boundaries, so the

reverse-engineering server is restarted between samples, which recovered 12 of 14 samples previously recorded as unmeasurable. Two audits were added for the same reason: diffing the claim register against the evidence log, and counter-searching each novelty claim in an adjacent field’s vocabulary, which demoted five including the one central to this paper. A counter-search is recorded whether or not it demotes anything, because a searched absence is only credible if the looking is documented.

2.6. Statistical analysis

Every arm is scored on the same samples in the same order, differences are computed per sample and aggregated, and 95% bootstrap confidence intervals (CIs) are taken over the paired deltas. The observations are nested, many claims coming from one sample and many samples from one family, so the bootstrap resamples whole clusters rather than rows and the p-value comes from flipping the sign of whole cluster means. Each interval carries the intraclass correlation coefficient (ICC) of its measure and the effective sample size that implies. Two properties of that test are reported with it. With k clusters the two-sided version cannot return a p below

$$p_{\min} = 2/2^k \quad (2)$$

At $k=5$ that is 0.0625, which puts $\alpha=0.05$ out of reach whatever the effect size. Above 20 clusters the 2^k assignments are sampled rather than enumerated, at 20,000 draws, with the observed assignment counted in its own reference set. The smallest p that route can return is therefore one over the draws rather than zero. Each reported p names which route produced it. The two properties pull in opposite directions: raising the cluster count lowers the floor and is what makes enumeration stop being possible, which is why the consensus arms were re-run one sample per family. Multiplicity is controlled with Benjamini-Hochberg over two declared families, and every null is reported with the minimum detectable effect of the design that produced it at 80% power, because a null without its resolution reads either as equivalence or as a failed experiment depending on the reader.

Both choices were corrections rather than plans: every interval in an earlier draft resampled rows rather than clusters, and an early study ranked prompt structures on a single-run parsed-claim count whose ranking inverted when the decoding budget changed. No result in this paper rests on a single run.

3. Results and discussion

3.1. The no-LLM baseline and the pipeline measured against it

Every accuracy number here is read against one reference: the sandbox the pipeline is built on, mapping its own signature hits to ATT&CK technique identifiers deterministically, with no language model anywhere. Table 3 reports it with the like-for-like comparison against the pipeline.

Table 3: The no-LLM baseline and the pipeline against the engine it is built on. Both blocks are scored per sample against family-level MITRE ground truth by the same code; intervals resample families, not samples.

	mean	95% cluster CI	ICC	effective n
<i>baseline alone, n=97 samples over 24 families</i>				
precision	0.2413	[+0.1850, +0.3010]	0.584	35.0
recall	0.1328	[+0.0896, +0.1810]	0.638	33.0
F1	0.1526	[+0.1114, +0.1998]	0.692	31.2
<i>paired, n=13 samples over 8 families, F1</i>				
CAPE alone, no language model	0.1130	[+0.0654, +0.1624]		
pipeline, static evidence only	0.1130	[+0.0770, +0.1543]		
pipeline, with the sandbox report	0.1160	[+0.0823, +0.1485]		

Ground truth is resolved per family, so two binaries of one family are scored against a byte-identical label vector and are not two observations: resampling samples rather than families gave an interval 2.32 times narrower at an intraclass correlation of 0.692, leaving 31.2 observations supporting an inference rather than 97. The sandbox asserted at least one technique on every sample, so this is a predictor rather than an artefact of an empty one.

The paired block is the 13 samples of the 97 that completed both arms, all three scored on one truth. On those, the full pipeline lands +0.0030 F1 from the signature engine it is built on, 95% cluster CI $[-0.0353, +0.0344]$, exact cluster sign-flip $p = 0.9609$, better on 7 of 13. The samples differ individually in both directions, so the system is not reproducing the sandbox’s answers but reaching different answers of the same quality. At 8 families the design could detect 0.085 F1 and did not, which bounds the pipeline’s contribution rather than showing it is zero. The cohort is 97 rather than 100 because 56 of the submitted batch never ran while the sandbox reported them complete, the fifth boundary case of Section 3.8. That paired block is not a random subsample of the 97, and the bound that produced it selects. Completion required the dynamic arm to finish inside its wall-clock ceiling, and the samples that did are the older ones: the paired block’s median year is 2021 against the cohort’s 2023, and the 26 samples dated 2025 and 2026 contribute none. The comparison therefore covers 8 of the 24 families and the early end of the cohort, and +0.0030 is a statement about that subpopulation rather than about the cohort. What the ceiling selects for beyond age is not measured, and the temporal study that could have said was withdrawn.

The three arms’ intervals overlap (Figure 2). For the equal-budget arms of Section 3.3 no baseline is definable at all: a regular expression over the dictionary that generated their evidence scores 1.000 on all 5 fixtures by construction, which is why those arms are drawn on their own axis. Put on one scale with these, a no-LLM baseline near 0.1526 beneath arms near 0.4136 would read as a large gain for the pipeline, and it is not one.




Figure 2: The one population where the pipeline and the baseline can be compared. Points are means over per-sample F1, bars are 95% intervals resampling the cluster the observations are independent at.

3.2. The deterministic technique assigner

The pipeline never asks the model for a technique identifier. The analyst describes behaviour and a retrieval index over the official ATT&CK corpus assigns the identifier, which removes identifier recall from a model that does not have the taxonomy memorised and cannot be made to by prompting. That assignment is the one component measured against external ground truth rather than our own cohort, on two public corpora neither of which is the ATT&CK text the index is built from (Table 4).

Table 4: The technique assigner on two external corpora: TRAM2 [11], sentence classification, and AnnoCTR [22], expert entity linking in running report prose. MRR is the mean reciprocal rank of the correct technique, gate separation is Equation (3), and AUROC is how often a correct top-1 outscores a wrong one.

backend	TRAM2, n=4,913				AnnoCTR, n=3,289			
	top-3	MRR	gate	AUROC	top-3	MRR	gate	AUROC
lexical	0.329	0.274	+0.085	0.690	0.183	0.147	+0.135	0.796
semantic	0.392	0.319	+0.019	0.606	0.214	0.176	+0.062	0.770
hybrid	0.392	0.319	+0.115	0.731	0.214	0.176	+0.168	0.818

Gate quality is the margin the absolute threshold has to work with, the mean retrieval score of the top-ranked technique when it is the right one less that score when it is not:

$$g = \text{mean}(s_1 \mid \text{correct}) - \text{mean}(s_1 \mid \text{wrong}) \quad (3)$$

Ranking quality and gate quality are separable axes and no backend wins both. The semantic index ranks better on both corpora while its scores barely separate a correct pick from a wrong one, because everything it retrieves sits near the same similarity. The lexical index ranks worse and gates cleanly, scoring near zero for unrelated evidence. Composing the per-axis winners matches the semantic index’s ranking

and produces the cleanest gate of the three on both corpora. Gate separation is a difference of means, so it compares backends only where their scores share a scale, and these do not: a correct pick scores 0.746 under the semantic index and 0.271 under the composition. The area under the receiver operating characteristic asks the scale-free question instead, how often a correct top-1 outscores a wrong one, and returns the same ordering on both corpora. That question could not be put to the first version of this evaluation, which averaged the per-item scores and kept only the means; they are retained now. That composition is the shipped default, and the gate is what makes the assigner safe to leave authoritative: an identifier it cannot align is refused, not guessed.

All three orderings replicate across the two corpora. Absolute accuracy does not and is not compared: entity linking in unconstrained prose over a 691-technique space is harder than sentence classification, and 48 AnnoCTR labels fall outside our ATT&CK bundle. The direction is not ours either, Büchel et al. [1] reporting the same ordering in a stronger and earlier form; what this adds is that the two axes come apart and that composing them beats choosing between them.

That layer also emits the artefact, and its conformance is measured with a standard external validator rather than with the integrity pass this project wrote itself: the OASIS cti-stix-validator 3.3.1 finds 0 errors and 0 referential warnings in a bundle built from the production models. Two defects the integrity pass had no opinion about were found that way and are fixed. Every object lacked the `spec_version` property STIX 2.1 requires on each STIX Domain Object rather than on the bundle, so no bundle the project had emitted was identifiable as STIX 2.1, and `null` and empty-array properties were serialised as present, which the specification forbids. The repair pass is not what makes a bundle conform: across 4 bundles with injected defects it recovers 0 to validator-clean, and what it does is preserve 53 objects that rejection would discard.

3.3. Multi-agent consensus at equal budget

Three arms at the equal budget of Section 2.5 run first on the fixtures and then again on real binaries, one per family so that every cluster holds a single observation (Table 5, Figure 3).

On the fixtures neither contrast is readable. The negotiated arm returns -0.0161 , $[-0.1247, +0.0819]$, at $3.2\times$ the tokens. That is two orders of magnitude below the design's 0.223. The noise control moves -0.0770 , $[-0.1337, -0.0233]$, with all 5 samples in the same direction, also below its own 0.120. Its exact permutation test returns $p = 0.0625$, the floor of Equation (2) at 5 clusters. An earlier draft read that separation as proof of sensitivity; what it shows is that all samples moved the same way.

The real-sample replication is where the design can see, the sign-flip floor at 24 clusters being 0.00005. Both multi-agent arms beat the single judge and neither beats the other: negotiated minus single is $+0.0537$, $[+0.0362, +0.0714]$, $q = 0.0001$, better on 19 of 24 families, while noise minus single is $+0.0532$, $[+0.0235, +0.0833]$, $q = 0.0040$. The contrast that isolates the mechanism returns $+0.0005$,




Figure 3: The equal-budget arms and the frontier comparison of Section 3.4, on the fixture corpus. This axis is not the axis of Figure 2: the two are different populations with different ground truth, and the fixture corpus has a ceiling a real sample does not.

Table 5: The equal-budget consensus arms, on synthesised fixtures and on real binaries. The real-sample block is scored against the same family-level ground truth as Table 3; the fixture block is not comparable to it.

arm	mean F1	rows	output tokens	× single
<i>5 synthesised fixtures</i>				
single judge, all evidence at once	0.4136	25		
negotiated multi-agent consensus	0.3975	25		3.2
noise control	0.3366	25		
<i>24 real samples, one per family</i>				
single judge	0.0967	1	39,986	1.00
negotiated consensus	0.1503	4	110,395	2.76
noise control	0.1498	4	94,388	2.36

$[-0.0294, +0.0312]$, $q = 0.9751$, on a design resolving 0.046 F1. The effect the negotiation would have to carry is the $+0.0537$ its neighbours show, which is above that resolution. If the negotiation were doing the work, this design would have seen it. What separates the multi-agent arms from the single judge is therefore what they share: 4 model calls instead of 1, and with them $2.76\times$ the output tokens. Spending more on a sample helps and making the calls argue does not, at this cap, on this corpus, with this model. What this design cannot do is separate the decomposition from the generation those extra calls bring, because only the two multi-agent arms are matched on calls; the contrast that is matched is the one returning $+0.0005$.

The two corpora gave opposite signs. On fixtures the noise control ran -0.0770

against the single judge and was worse on every sample; on real binaries it runs +0.0532 and is better, at the same q as the negotiation. Nothing about the arms changed and the evidence did. A corpus whose inputs are generated from its own answer key does not merely have less power than one of real binaries; it can rank the arms the other way round.

3.4. Model scale, quantisation and the reasoning flag

Table 6 reports three comparisons at the same 2,400-token output budget: a 3.4 $\times$ larger frontier model, and the same weights hosted at full precision with and without the reasoning flag.

Table 6: Model scale, quantisation and the reasoning flag, all paired against the local arm on the same 5 fixtures and repeats. The local arm is Qwen3.6-35B-A3B at IQ3_K_R4, the frontier arm Nemotron-3-Super-120B-A12B, and the vendor-hosted arms the same Qwen weights at full precision.

arm	think	F1	vs local	95% CI
local, 3-bit	off	0.4136		
frontier, 3.4 $\times$ larger	on	0.4162	+0.0026	[−0.1322, +0.1460]
frontier, flag requested	on	0.4149	−0.0014	[−0.0695, +0.0799]
vendor-hosted, full	off	0.3507	−0.0629	[−0.2133, +0.0963]
vendor-hosted, full	on	0.0080	−0.4056	[−0.5232, −0.2880]

Two of these three comparisons are bounds the design cannot read, and are reported as such. The frontier model lands within three thousandths of F1 of a 35B model quantised to run on one desktop graphics processing unit. It is better on 12 and worse on 13 of 25 rows. The design’s minimum detectable effect is 0.300 F1, coarser than every architectural effect reported here. It is confounded as well: the local arm runs with its reasoning stream disabled, necessarily, because the local server otherwise returns an empty answer, while the frontier arm spent 56.5% of its budget reasoning, and re-running it with the parameter set returned 56.2%. The arm was not a collapsed one, having parsed 25 of 25 calls with no refusals and 1 hitting the output cap; an earlier version of it reported 0.5025 at $n=9$ and suggested a lead that did not survive completing the sample. We report the null because the literature predicts otherwise: parameter size is the only statistically significant predictor of ATT&CK-classification F1 on the nearest task [47] [preprint]. Nor can we test it as a series. A rank correlation over model size needs every arm on the same side of the reasoning flag, which is worth more than the size effect, and the only endpoint above 35B does not let that flag be set. Quantisation is the second bound: the paired difference is −0.0629 on [−0.2133, +0.0963] at $n=25$, admitting a penalty of up to 0.2133 F1 against a minimum detectable effect of 0.331, so 3-bit quantisation is not demonstrably this pipeline’s handicap rather than not one.

The third comparison returns the largest effect measured anywhere in this paper, and its sign is the one that has to be stated carefully: what gains is turning the reasoning flag off. It replicates on two endpoints. On the model we host, disabling it gains +0.3427 paired ([+0.2008, +0.4936]), with 24 of 25 calls exhausting the

output budget on reasoning. On a third endpoint, enabling it spent 99.9% of every call’s budget thinking and scored 0.0000 across 25 calls, against 0.4501 with it disabled: a difference of +0.4501 F1, [+0.3604, +0.5437]. The flag is the only effect in this paper that exceeds its own design’s minimum detectable effect, at 0.310 here and 0.200 on the third endpoint: the largest effect measured on any arm is a configuration flag, not an architecture and not a parameter count. The result is exploratory rather than confirmatory, and nothing here should be read as a test of a declared hypothesis. No prediction about the flag existed before the confound was found; the comparisons were run afterwards, and corrected as a family of 5 their q-values are 0.1042 and 0.1042, neither of which clears $\alpha=0.05$. Like every comparison on this corpus they also sit at the exact test’s floor of 0.0625. What carries the finding is the size of the effect against the design’s resolution, not its q-value.

3.5. Retrieval, confidence and firing rate before effect

Three retrieval components were built independently, each works where it was built, and each moves nothing where it is wired in (Table 7).

Table 7: Three retrieval components, each measured in isolation and again end to end.

component	measured in isolation	contribution end to end
ATT&CK case-prior RAG	self-consistency 0.620 against a 0.424 frequency prior, scored on the corpus’s own prior attributions rather than on MITRE ground truth	loses to the frequency prior in production (0.111 against 0.123, MITRE ground truth)
family-feature RAG	yes, recall@5 0.199, 6.3× chance, leakage-free split	+0.0029 F1, precision −0.0088, n=19
opcode-hash attribution	n/a, never fires	0 of 18 samples; 7,716 functions hashed, 0 matches

The first row’s isolation figure is not accuracy, the leave-one-out corpus being labelled with the pipeline’s own prior attributions; it is kept because a frequency prior scored the same way reaches 0.424 and the production column beside it is scored against MITRE ground truth. The three fail differently. The first fails on vocabulary mismatch between query and corpus, visible only in the score distribution while top-k accuracy stayed healthy; the second is simply small; the third fails twice over, its corpus holding 3 samples under 2 generic labels while the 8-instruction floor that correctly excludes thinks leaves 9 of 18 real samples with one hashable function or none.

Verbal confidence, which the cascade and every deterministic gate consume, discriminates correct from incorrect claims at an area under the curve of 0.550, near chance, concentrated on 4 round values. Resampling samples rather than claims puts the interval at [0.475, 0.671], which contains chance, and 3 of the 5 samples sit at or below chance individually, with the pooled figure carried by the one reaching 0.778. Figure 5 shows why: 186 of the 210 claims carry confidence exactly




Figure 4: Firing rate decides whether an ablation can be read at all.




Figure 5: The number the gates consume barely orders anything. a) the receiver operating characteristic over 210 claims against resolved ground truth, b) the confidence actually emitted.

1.0, so the score has almost no ordering to contribute. The estimate is rank-based with ties averaged, and sorting them arbitrarily returns 0.458 instead, a difference entirely an artefact of how 89% of the data is handled. This matches Kumaran’s finding [54] [preprint]: verbal confidence measures willingness to commit rather than correctness, so every deterministic gate keyed to that number is keyed to something we cannot distinguish from noise.

Whether any of these mechanisms can be ablated at all depends on how often it engages, which Table 8 and Figure 4 report. Three of the four mechanisms the architecture was built around engage on almost nothing, so an ablation of any would return a null describing the cases where the mechanism never ran. The cap’s pre-conditions explain its rate: it applies to three technique families only when the sole contributing layer is static, and 84% of those claims are YARA-only. Only the sink-

reachability hint clears a rate at which an ablation carries information, which is why it is the one we ablated.

Table 8: Firing rate before effect, and the one ablation it justified. Units differ per row of the lower block and are named in the row.

mechanism	fires on	consequence
confidence cap	0.82% of techniques	an ablation would measure nothing
sink-reachability priority hint	56.7% of samples (55/97)	an ablation is interpretable and was run
STIX integrity pass	60 fresh judge bundles	reported by removal reason
outcome of the hint ablation	mean Δ (on – off)	95% CI
distinct technique identifiers	+0.50	[−3.33, +4.50]
claims	−0.83	[−4.33, +3.67]
seconds	+52.55	[−161.93, +268.45]

The hint does nothing measurable. On $n=6$ pairs it is better on 2, worse on 2 and tied on 2, with per-pair deltas of +9, +4, 0, 0, −4, −6, large in both directions and cancelling. This is the fourth architectural claim to end in a null and the last one we were in a position to test.

That null is constrained by a second result: half the experiment was not scoreable. Twelve pairs were attempted and 6 survived. 3 were lost to degenerate decode, an arm emitting 49 to 117 claims across only 2 to 14 technique identifiers; 2 were unattributable, both arms dead with the host state never recorded; and the last was incomplete at the 630 s bound, a genuine outcome that still costs the pair. A 50% pair-loss rate bounds what any ablation here can detect, and quoting +0.50 [−3.33, +4.50] without it would claim a precision the instrument does not have. The unattributable pairs are this paper’s own retention rule turned against it: we had retained per-sample outputs, and what went unrecorded was the host state the arms died under (Section 3.11). Two further arms carried a `Connection error` from a restart race, meaning the measurement never happened, and were re-run, contributing a −4 against the hint; a third was the timeout, which is an outcome, and deleting an outcome one dislikes is selection rather than repair.

Running the ablation also surfaced M4 of Section 3.8, a production defect that 2,716 passing tests did not: on any binary rich enough to exhaust the 40-step Re-Act budget the analyst returned zero techniques. The fix was verified on the same sample, which went from 1,677 s to 323 s and from 0 to 5 technique identifiers. It is partial. One sample still exceeds the bound, and generation on this hybrid recurrent architecture varies from 162 to 20 tokens/s (Table 1), so bounding the input does not rescue a case where the tokens arrive eight times slower than the budget assumed.

3.6. The reach of the corroboration cascade

We remove one Layer-0 source at a time and compare the bundle the analyst receives against the one produced with all of them, under two conditions differing in whether the removed source solely owned its techniques. Under *disjoint*, where its techniques disappear because nothing else claims them, the verdict changes in 32 of 32 arms at Jaccard 0.750 [0.714, 0.800]. Under *overlap*, where the techniques survive under a partner source but their corroboration changes, it changes in 0 of 32 at Jaccard 1.000 [1.000, 1.000]. Removing a technique removes it from the bundle; destroying the cross-domain agreement behind it, which is what the cascade exists to compute, leaves the bundle byte-identical. Perturbing the trust weights reaches no further. Across 5 perturbations on 97 samples, one of which inverts the most- and least-trusted layers, the top-10 ranking changed on 12.4% to 28.9% of samples and the corroborated set changed on none, because corroboration depends on how many layers agree rather than on how far each is trusted.

The reason is not the one we first wrote, and the correction is M5 of Section 3.8. The reconciliation step does two things. It drops every attack-pattern whose technique identifier cannot be resolved, and it appends every cascade technique missing from the bundle. The cascade’s set is therefore a guaranteed subset of every bundle the system emits, whatever the judge produced. Recomputing each arm’s cascade set from the same seeded fixtures shows the bundle is exactly that set in 80 of 80 arms. The judge adds no technique the cascade did not already hold: it emits attack-patterns, but never one that is new. Both conditions therefore follow from set arithmetic rather than from restraint. The Jaccard value was the signal we missed, for the reason M5 gives.

This ablation replaces an earlier version that reported the verdict unchanged in 0 of 15 cases. That version carried two defects. The sources it varied excluded the Sigma layer, which fires on 94 of 97 samples at weight 0.55, above every static layer it did vary. And at five techniques over six sources, the round-robin assignment left one source with nothing to remove, so its arm matched the baseline by arithmetic rather than by result. The re-run uses fixtures of 12 to 51 techniques so every source carries at least three claims, includes Sigma, and adds a corroborated pair that does not involve YARA; the null survives at 32 arms rather than 15. One pre-registered prediction resolved in both directions: removing the tool-artifact layer should change nothing, because it emits on YARA’s cascade domain and cannot contribute corroboration, and under overlap that is right at 0 of 8 while under disjoint it is wrong at 8 of 8, because there the source solely owns its techniques. The prediction holds precisely where corroboration is the only thing varying.

3.7. What the judge contributes

The set arithmetic leaves two pipelines the evidence cannot tell apart: a judge whose output is unusable and wholly replaced, and one that reproduces the cascade’s set exactly. Intercepting the reconciliation step and recording what the model produced before the deterministic set was merged in settles it (Table 9): three of the

four calls produced nothing nameable, and the fourth named only techniques the cascade already held. The bundle carries the judge’s name and the cascade’s contents.

Table 9: What the judge contributed to the artefact the analyst receives, intercepted at the reconciliation seam. Counts are techniques, over four calls.

	total	per call
attack-patterns the judge emitted	50	12.5
carrying a resolvable ATT&CK identifier	12	3.0
dropped, the model named no technique	38 (76.0%)	9.5
identifiers the cascade did not already hold	0	0.0
injected because the judge omitted them	87	21.8

Half the calls never reached that step. Four of the eight timed out at the verdict ceiling, and on a timeout the verdict tool returns a text-fallback bundle that never calls the reconciliation routine, so the cascade is not consulted at all. The timeouts are not a property of the fixtures but of M7 (Section 3.8): the judge’s 8,192-token output cap reached the local server under a name it does not read, so the model decoded unbounded until the caller gave up, one such call measured at 30,155 generated tokens and still going.

Repeating the study with the cap fixed turns the pair into a controlled experiment, and the result inverts this section’s own conclusion. With the ceiling binding, the same four fixtures fail and the same four succeed, the judge’s share of the bundle is 0.0% in both conditions, and every number on the reconciled path is identical. Only the way the failure arrives changes. The artefact does not. The fallback builder scraped ATT&CK identifiers from the evidence claims and from the model’s own raw response; in the uncapped condition that response is the literal string [TIME-OUT] and contains none, so the bundle was exactly the evidence-derived set, identical to the cascade set on all four calls. In the capped condition it is 18,748 to 24,710 characters of degenerate decode (Table 10).

Table 10: Techniques reaching the analyst on the text-fallback path against those the corroboration cascade held, per fixture. The last column is the set no evidence source corroborated.

fixture	fallback	cascade	only in fallback
jhuugit	32	20	12
sardonic	27	25	2
sliver	46	23	23
wannacry	26	16	10

47 techniques reach the analyst that the corroboration cascade never held, and none goes the other way; on one sample the bundle doubles. Because the uncapped run is a control whose raw text provably contains no identifiers, every one is attributable to the model’s own output, and they passed through no filter at all: not

the cascade, not the reconciliation step, not the invalid-identifier filter, not the integrity pass. Checked against the 691-technique catalogue, they are also real: 45 of 45 unique identifiers exist and none is invented. That removes the easy reading. These are not hallucinated identifiers a schema check would catch. They are genuine ATT&CK techniques that no evidence source claimed, carried in the same object type and the same bundle as techniques three independent sources agreed on, with nothing downstream to tell the two apart. The judge therefore has no influence over which techniques reach the analyst on the path where it works, and more influence on the path where it fails: 0 of 99 techniques when reconciliation runs, against 47 of 131 when it does not. The architecture suppresses the model's contribution precisely when the model is functioning, and admits it precisely when it is not.

That path is a fail-open one and it is now closed. The fallback builder no longer promotes an identifier that appears in the model's raw response and in no evidence claim; it records what it declined and emits only what an analyst claimed against a cited artefact, so the degraded path emits less than the working one rather than more. The measurement above describes the pipeline as it ran, and is reported as measured for the same reason the other repairs in this paper are: a fix applied before the number is taken is a number about the fix.

3.8. Instrument failures

Seven failures of the instrument share a shape: it answered successfully, and answered about something else. Table 11 lists them and Figure 1 marks each on the topology. The first four crossed a boundary with another server. The last three are ours, two in the analysis code that produces our results and one in the pipeline that produces the artefact the analyst receives, so the shape stops neither at the network boundary nor at the evaluation harness.

M2 is the mechanism the others are read against, and what makes it one is that the response contained the correct program name. The server keeps its own notion of which program is current, and `load_program` sets that only when nothing is current yet:

```
load A      -> {"success": true, "program": "A"}   current: A
load B      -> {"success": true, "program": "B"}   current: A   <-- still A
run_analysis -> {"program": "A", "new_functions": 0}
call graph  -> A's graph, byte-identical, for every sample
```

From the second sample of a container's lifetime onwards, the decompilation, the imports, the strings, the call graph and the function hashes all described the first binary while the report named the current one. M1 and M3 are the same shape at neighbouring layers: M1 survived to first contact because the reverse-engineering server declares no optional parameter with a typed default and was the only server that client had talked to, and M3 is error swallowing in the ordinary sense, `raise_for_status()` seeing nothing wrong in a body that says the load failed. M4 composed two bounds each correct alone.

A fifth case at the same boundary was caught before it produced data, and it is why the cohort is 97. A request for a report the sandbox has since deleted is

Table 11: The seven instrument failures, with the layer each occurred at and what found it.

	what the instrument did	layer	what found it
M1	an optional argument the agent never supplied arrived as an explicit <code>null</code> , and a field typed <code>str</code> rejected it, failing all 36 tools	sandbox MCP client	every call returning the same validation error
M2	<code>load_program</code> answered success without changing the server’s current program, so everything derived from it described the first binary of that container’s lifetime	reverse-engineering server session state	call graphs identical to the character across unrelated samples
M3	a load the server could not perform answered HTTP 200 with an error in the body, so downstream code analysed whichever program remained current	HTTP status against body	66 consecutive samples at exactly 75,426 characters
M4	a 40-step ReAct budget and a wall-clock cap composed into a path that returned zero techniques, at 28 minutes per attempt	two local bounds	a 25-minute call always producing one claim and zero techniques
M5	an ablation varied a deterministic reconciliation step and the output was attributed to a language model	our analysis code	a Jaccard too clean for a model to produce
M6	a rank correlation over model size ranked a reasoning-configuration flag instead, recovering the published scaling coefficient	our analysis code	reading the arms table under the correlation
M7	a documented output ceiling was renamed during serialisation to a key the server does not read, so a component described as bounded was never bounded	our production pipeline	a study recording which failure branch each call took

answered with HTTP 200 and a short body saying the reports directory does not exist. A fetcher trusting the status code would have written those bodies into the archive under the filename of the sample they were meant to describe; ours verifies the report’s own digest against the identifier it asked for before writing, a four-line check that exists only because M3 taught us the shape. We first recorded the consequence as a retention limit, that the reports had existed and expired, and that was wrong. Asking the sandbox how long each analysis had taken produced a split with nothing between the modes: the 43 tasks whose reports survived ran 186 to 366 s, and the other 56 ran under a second, all 56 of them still marked reported. That accounts for 99 of the 100; the last returned no duration at all. A Windows PE does

not detonate in one second. Nothing had expired; nothing had happened, and the instrument said otherwise. Re-submitting recovered 54 of them, and the remaining 3 produced no report on either attempt, all 3 still marked reported at zero seconds. A completion status is a statement about a queue rather than about an analysis, and is exactly as trustworthy as an HTTP 200; the retrieval path now refuses any task under 60 s.

The last three crossed no boundary. M5 is the only one that reached a results table. We ablated the corroboration cascade and wrote the two conditions of Section 3.6 up as a finding about the judge, when the reconciliation step makes the cascade's set a guaranteed subset of every bundle whatever the model produced, so both numbers are necessities of an `if` statement and the experiment could not have returned anything else. A second study about to run would have compared a weighted cascade against a flat union of the same claims, sharing one reconciliation set, and was stopped four minutes in by the log line that exposed M5. The signal was in the results the whole time: a Jaccard of 1.000 with a zero-width bootstrap interval across 32 ablated arms at temperature 0.1 is not something a language model produces. It is what a constant looks like. We read it as an unusually clean null, because a clean null was the result we were prepared to find.

M6 is the only one whose wrong answer agreed with the literature. Assembling a parameter-size series, the harness ranked mean F1 against total parameters across a $3.4\times$ span and reported $\rho=+0.866$, reproducing the prior almost exactly, with the confounds it could not remove printed honestly beneath. The five rows were three models: one appeared twice because it had been run in two configurations, and the same weights score 0.3507 with the reasoning stream disabled and 0.0080 with it enabled, because 24 of 25 calls then spend their whole output budget reasoning and never answer. The enabled row sat at the small end of the parameter axis, which is where it set the sign of the correlation. The failure is not the duplicate rows but that the arm-selection rule did not exist: the harness had averaged ranks, an exact permutation p and a common-cell restriction, every guard about a way the numbers could mislead and none about whether the rows were comparable. M5 was caught by a number too clean to believe; M6 produced one in exactly the range a reader would expect, so nothing was anomalous to notice, and it was caught by asking why one model was on the arms table twice.

M7 is in the production pipeline and removed a guarantee the code states in its own comment. The judge is built with an 8,192-token output ceiling to bound verdict generation, the client library renames that parameter to the vendor's newer spelling when it serialises the request, and the local server does not read the newer spelling: it accepted the field, ignored it, and decoded without a ceiling, 30,155 tokens generated and still going when the caller's ten-minute wrapper gave up. The consequence is the four unreturned verdicts of Section 3.7. It was invisible because the configured value was correct at every level anything reads, and only the serialised request was wrong, which nothing in the system reads. What found it was a study measuring something else, which recorded which branch each failed call took. The incompatibility is known, reported against `llama.cpp` in its own issue

tracker [55] [preprint]; what one costs inside a measurement is the part we report.

That 2,716 tests passed throughout is a fact about test shape rather than test quality. M2 and M3 need a second case within one server lifetime, and a unit test loads one program, asserts and tears down; M1 needed a second server; M4 needed an input rich enough to exhaust the step budget. Preconditions of that kind are what a fast unit suite is designed to avoid. The last three are invisible to unit testing in principle, because every unit behaved as specified: M5’s reconciliation function was never wrong and no test of one component catches a misattribution made three modules away, every one of M6’s guards concerned the wrong question, and a unit test of M7 would check that the ceiling is 8,192, which it is, at every level the object model exposes.

3.9. Output cardinality as a detector

Five of the seven were found by the same observation, a number that repeated where variation was expected (Table 12). M6 and M7 were not, because their outputs vary and there is no repetition to notice.

Table 12: What the output-cardinality check saw and which defect each observation turned out to be.

observation	defect
a priority hint of exactly 2,575 characters for two unrelated binaries, and a third in an earlier session	M2
call graphs identical to the character (404,337 chars, 11,798 lines) for samples of 241 KB and 139 KB	M2
66 consecutive samples at exactly 75,426 characters	M3
every tool call returning the same validation error	M1
a 25-minute call that always produced one claim and zero techniques	M4
a Jaccard of 1.000 with a zero-width interval across 32 ablated arms at temperature 0.1	M5

The generalisation is one line of arithmetic. Before a batch measurement is trusted, count how many distinct outputs the N inputs produced, on a dimension that should vary, such as output length, digest or element count. If far fewer than N of them differ, the instrument is repeating itself, and repetition is what stale state looks like from outside. Plotted against inputs processed, distinct outputs track the diagonal for a healthy instrument and go flat for a stuck one, so the diagnosis is in the shape (Figure 6). M5 extends it to our own results.

The price is a completed study. A seven-year 210-sample temporal-drift analysis ran the full pipeline per sample against a long-lived server never restarted between samples, which is M2’s precondition exactly, and whether it was affected cannot be determined because its per-sample outputs were not retained. The withdrawal is caused by the defect plus the inability to re-ask the question afterwards.




Figure 6: The whole detector, drawn. a) a measured run at 63 distinct call-graph sizes across 97 samples, b) a schematic of a stuck instrument. Panel b is drawn rather than measured because the run it depicts predates the per-sample retention these failures caused us to adopt, so its sizes no longer exist to plot.

3.10. Discussion

Most of these negative results are bounds rather than equivalences, which was not evident when they were first written down: five of the studies run on five fixtures, where the sign-flip floor of Equation (2) puts $\alpha=0.05$ out of reach whatever the effect size. That also inverts the cheap fix. On an underpowered paired design the instinct is to add repeats, and on clustered data a sixth fixture would have halved the attainable p while a sixth repeat on five fixtures moves it not at all.

The literature we counter-searched against locates silent failures at the boundary between modules, and four of our seven crossed such a boundary while three crossed none. What the seven share is not a boundary but an attribution: an ablation varied a deterministic reconciliation step and we attributed the difference to a language model, a rank correlation ordered a configuration flag and we attributed it to parameter count, and a renamed output ceiling let us attribute boundedness to a component that had never been bounded. The number was real each time, and the arrow pointing to its cause was drawn by us. Seven cases from one project is not a class, so this is a proposal rather than a taxonomy, but it suggests a different question from the one boundary-oriented thinking asks: not whether a component failed silently, but whether each number's cause was measured or assigned. The clustered-inference correction is that question turned on ourselves, answered badly.

None of this shows negotiation to be worthless, model capacity to be off the ceiling, or verbal confidence to be uninformative. Each holds at this design's resolution, on this corpus. What we can say is what each mechanism cost.

3.11. Limitations and threats to validity

This section is organised around the nine pitfalls of *Chasing Shadows* [37], with two checks added that the survey does not have and that both caught errors here. It begins with damage rather than with hazards. The withdrawn drift study of Section 3.9 had entered our claim ledger as a novel result, and it is not restated anywhere in this paper. Table 13 gives our exposure to each pitfall with the basis for the verdict.

Table 13: Our exposure to each of the nine pitfalls, with the basis for the verdict.

pitfall	verdict	basis
P1 data poisoning	CLEAR	nothing is trained; retrieval corpora are curated and their provenance documented
P2 label inaccuracy	CLEAR	ground truth is MITRE-curated and no LLM scored any result; no human analyst rated any report either
P3 data leakage	PARTIAL	one instance found by us, disclosed and mitigated, but never systematically audited; the model’s training data is unknown to us and memorisation was not probed
P4 model collapse	CLEAR	augmenting the long-term memory corpus with LLM-generated cases was considered and refused
P5 spurious correlations	PARTIAL	two were found and are reported; no systematic perturbation testing was performed
P6 context truncation	EXPOSED	truncation is designed in throughout and the counters now exist, but the distribution over a full cohort does not
P7 prompt sensitivity	PARTIAL	a structured prompt-variation study exists, but production prompts are fixed and were not varied per model
P8 surrogate fallacy	PARTIAL	every local result comes from one model on one machine, and a second endpoint has now run to completion against it
P9 model ambiguity	PARTIAL	model and engine are pinned by digest, revision and commit, but the commit was reconstructed rather than recorded

Three of the nine need more than a row. Truncation (P6) enters through chunking at function boundaries, a per-call evidence cap, a 40-step ReAct budget, an 8,192-token verdict cap and a 400-char hint cap, and that enumeration was itself incomplete: a 20,000-edge cap on the call-graph fetch was found only while instrumenting a different measurement, and silently truncates 1 of 97 samples. Bounds also compose, as M4 shows, and removing a hint made the judge overrun a 600 s ceiling and return an empty bundle 6/17 times.

P8 is the pitfall this work is most exposed to and the one where the evidence has moved. A 120B reasoning model ran to completion against the local model and did not separate from it, and quantisation is no longer unmeasured either (Section 3.4): a 3.4× parameter advantage producing no measurable difference is evidence against

the pitfall’s usual worry rather than an absence of evidence for it. Both bounds are confounded, the null by a reasoning flag that cannot be removed on that endpoint and outweighs the size difference, the quantisation figure by two serving stacks that differ alongside the precision. What P8 cannot close is the size series, for a measured rather than a budgetary reason: a parameter-size trend needs three arms at three sizes on the same side of the flag that outweighs size, the one endpoint honouring that flag hosts the 35B model we already run, and the one above 35B does not honour it. Coverage remains as well, the second model having been measured on fixtures rather than on the malware cohort. P9 is PARTIAL because the running binary reports its version as unknown and the engine commit was reconstructed rather than recorded (Section 2.4): build provenance must be captured at build time, and ours was not.

Two checks the survey does not include, both described in Section 2, caught errors one level below what the nine pitfalls address: a component scored against labels drawn from its own source is scoring a tautology, and a retrieval corpus overlapping an evaluation cohort is scoring its own training set. Neither was assumed. Both were tested component by component, and the working is in the released artefact [53]. Two construct problems survive and both are ours: the ATT&CK case corpus is labelled with the technique identifiers the pipeline itself previously attributed, which is why that number is relabelled as self-consistency in Section 3.5, and the catalogue defining valid identifiers is the same file that bounds the label universe.

Any claim about what the sandbox contributes has to survive one measurement of what the sandbox reports. Across the 97 archived analyses there are 143 distinct network domains, of which 38 appear in every sample: the analysis virtual machine’s own vendor telemetry, present whatever was detonated. Per sample that share runs from 55.9% to 79.2%, median 67.9%. Two-thirds of the network evidence in a dynamic arm is therefore the instrument describing itself, and an effect attributed to malware behaviour may be partly that machine’s idle traffic. The same measurement bounds two Layer-0 sources from the other side: the living-off-the-land binary and domain generation algorithm layers produced a claim on none of the 95 reports archived when that ablation ran, while being fed a median of 8,667 recorded API calls and 48–68 domains respectively, so neither appears as an arm in the cascade ablation. An earlier version of this measurement, on the 43 analyses that survived before the cohort was recovered, gave 130 domains with 40 ubiquitous at a median share of 71.4%: more than doubling the cohort moved the figure by three points.

The machine a measurement runs on is part of the instrument, a threat we did not anticipate and met late. The working set does not fit a 31 GiB host alongside an interactive session, with the model server holding ~16 GB, the analysis worker ~8 GB and the disassembly container up to 6 GB. During a paired ablation the swap file was exhausted and the model server had 2.3 GB of its own address space paged out; an arm then exceeded a 594-second budget on a prompt trimmed to 16,000 characters. Whether it failed because of the pipeline or because of the host

cannot be determined, because we recorded the pipeline’s outputs and not the host’s state, so those arms are reported as `unattributable` and the ablation as `incomplete` rather than as a null. A wall-clock bound is therefore a host-dependent measurement presented as a system property, and any result that turns on a timeout is a statement about this pipeline on this machine under whatever load it had. The screen we added applies forward only: arms already collected without host state cannot be re-screened.

Three objections remain that the work does not close, and are argued in the artefact [53]: the judge findings rest on four calls rather than a rate, the detector’s own evidence is retrospective, and seven defects in one’s own code is a post-mortem rather than a result. We claim the setting and the price rather than the category, every mechanism having been counter-searched in an adjacent field’s vocabulary (Section 2.5).

All three are blocked by the same thing, a host on which the model server does not fit alongside anything else, so a long run walks the host into its own memory guard. Cycling the server from the checkpoint between stretches of work is what closed a fourth objection, that the equal-budget arms ran only on 5 synthesised fixtures whose evidence is in bijection with their answer key: two of the three were re-run on real binaries and the answer changed sign (Section 3.3). The frontier comparison and the confidence calibration still rest on that corpus and are reported as a pilot.

4. Conclusion

This work addressed two barriers to trusting a measurement of an LLM malware-analysis pipeline: that such pipelines are rarely scored against the deterministic tooling they are built on, and that the instrument producing the score is itself a distributed system whose failures do not announce themselves. The contribution is a way of reading what a component adds that does not depend on the component working, in four parts: a baseline to refer an F1 to, a firing rate to decide in advance whether an ablation can say anything, a minimum detectable effect to turn a null into a bound, and a count of distinct outputs to ask whether the instrument produced N answers or one answer N times. None is expensive and each changed a conclusion here.

Four architectural claims motivated the design and one survives in part. Where a corpus large enough to produce a significant answer was available, decomposing the analysis into several model calls scores +0.0537 F1 above a single judge at 2.76 times its output, while the negotiation between them returns +0.0005 at matched calls. The design helps, though not for the reason it was built. The corroboration cascade never reaches the artefact the analyst receives, two-tier attribution fires too rarely to carry the result, and the confidence number every deterministic gate consumes is indistinguishable from chance. What moves the end-to-end result is the layer underneath, where the deterministic assigner gives the cleanest identifier gate of three backends on two external corpora. The finding we did not expect to be

the paper is that the instrument was repeatedly wrong in ways that looked like data: seven mechanisms are catalogued with the layer each occurred at, four of them found by that same count of distinct outputs.

We are careful about what is new. Silent failure at tool boundaries is documented, the metamorphic relation behind the detector is ordinary, and inference-time configuration inverting model rankings inside a constrained budget is established. What we add is the setting and the price: one completed study was withdrawn because its per-sample outputs were not retained, and four arms of a later ablation were then lost the same way, one level up, because the rule written after the first failure was scoped to that failure.

This is one pipeline, one 35B model at one quantisation, on one machine, evaluated on Windows PE malware against family-level ground truth, and most of its nulls are bounds set by a five-cluster corpus. Future work follows from each: the withdrawn drift study re-run under the corrected instrument, the judge intercept extended from four calls to the full cohort, the parameter-size series on endpoints that honour the reasoning flag, and the output-cardinality check tested prospectively. Each needs a host the model server does not fill on its own, which is the binding constraint on this work rather than the method.

What we do assert is a change of default. In a pipeline assembled from other people’s servers the correctness of the model is not the binding constraint on the correctness of the result. Long before a benchmark score is a claim about a system, it is a claim about an instrument, and a server that answers is not the same as a server that answered the question it was asked.

Declarations

CRedit authorship contribution statement

Düzgün Küçük: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Data curation, Visualization, Writing (original draft). *Fatih Ertam:* Conceptualization, Supervision, Project administration, Writing (review and editing).

Ethics, containment and data sources

Every sample analysed in this work is live Windows PE malware. The 97-sample cohort and the 210 dated samples of the withdrawn drift study were obtained from MalwareBazaar [56], under terms that permit research redistribution of samples but not of the corpus as a package. MABEL, the Malware Analysis Benchmark for AI/ML, was mined offline into the family-fingerprint catalogue and the ATT&CK case corpus with no binaries downloaded, and Ultimate-RAT-Collection supplied the remote access trojan builders and payload binaries for that catalogue and for the leakage-free retrieval split. The ATT&CK Enterprise knowledge base [57] supplies the family-to-technique uses relationships used as per-sample ground truth, under the ATT&CK Terms of Use.

Detonation happens only inside a CAPE sandbox on a dedicated Windows guest image reverted to a clean snapshot between analyses. The sandbox host is reachable from one network segment and from no public route, holds no credentials for any other system, and shares no filesystem with the host that runs the pipeline. Binaries are held only in a directory excluded from the host’s real-time scanner, are never committed to version control, and are not redistributed with this paper. We note a limitation of that arrangement rather than claiming it sufficient: the guest reaches the public internet during detonation, because a sample that cannot resolve its command-and-control infrastructure behaves differently from one that can and the network evidence channel would otherwise be empty for every sample. Live detonation with egress is the standard arrangement for behavioural sandboxing, and it is the arrangement whose consequences we report. No user study was conducted, no personal data was collected, and no human analyst scored any output, that last point being a limitation of the evaluation rather than a claim about ethics.

Data availability and reproducibility

The complete project is released under an open licence as an archived software release [53]: the pipeline itself, the evaluation harnesses, the derivation of every number in this paper, the per-sample records they read and the scripts that build the figures and assemble the document. That release is what the statements below refer to, rather than the development repository, because a repository moves and a claim about an artefact should name something that does not. The archive holds every per-sample record retained by every study reported here, in the JavaScript Object Notation files the harness wrote, together with the offline re-analysis that recomputes every interval, design effect, minimum detectable effect and adjusted p-value from those records and its random seed, the ATT&CK ground-truth fixtures, and the model digest and engine commit needed to reproduce the local arm. Two runs of the offline analysis produce byte-identical output, asserted by a test in the released suite rather than claimed here, and every interval carries the seed, the iteration count and the number of clusters it was computed from. The live arms are not bit-reproducible and we do not claim they are: decoding is at temperature 0.1 on the local arm and at each provider’s default on the hosted ones, one of which accepts a reasoning-configuration parameter and does not act on it, so anyone re-running them should expect the numbers to move and should read the intervals accordingly.

Four things are withheld. Malware binaries are not released and are obtainable from the sources named above by the SHA-256 digests in the archive. Sandbox reports are not released in full because they embed strings and memory contents from the samples. The sandbox host’s address appears nowhere in the archive. And one dataset used for retrieval-corpus construction is redistributed by its authors under terms that do not permit our redistribution, so it is referenced rather than included. One thing is not available at all: the 210-sample temporal-drift study’s per-sample outputs do not exist, having been written to a path that was valid on the machine the harness was first written on and silently relative on the machine it ran on. The study is withdrawn on that account and the withdrawal is reported in the results rather

than quietly repaired. The manifest of the 210 samples survives and is included, so the cohort can be reconstructed even though the results cannot.

Use of generative artificial intelligence

Large language models are the object of study in this work, so the pipeline under evaluation is built from them and every measurement reported here is a measurement of their behaviour in it. Separately, generative artificial intelligence was used as an authoring and engineering aid, for drafting and revising prose, writing evaluation harnesses and their tests, and performing literature searches whose results were then verified by fetching each source. Of the nine citations added by those searches and checked one by one, six said something the source did not, and all six are corrected in the text, because a search summary that reads as a citation is the same class of failure this paper is about. An internal readability assessment of report narratives was performed with a language model as a development aid; its results are deliberately excluded and no LLM scored any result reported here. The authors take responsibility for all content, including every number, every citation, and every claim about what the measurements support.

Declaration of competing interest

The authors declare no competing financial or non-financial interests, and no external funding supported this work. The three commercially hosted inference endpoints were paid for at their public list prices, and their providers had no involvement in, or advance knowledge of, this evaluation.

References

- [1] M. Büchel, T. Paladini, S. Longari, M. Carminati, S. Zanero, et al., SoK: Automated TTP extraction from CTI reports, are we there yet?, in: 34th USENIX Security Symposium, 2025, pp. 4621–4641.
- [2] H. S. Gunawi, C. Rubio-González, A. C. Arpaci-Dusseau, R. H. Arpaci-Dusseau, B. Liblit, EIO: Error handling is occasionally correct, in: 6th USENIX Conference on File and Storage Technologies (FAST), 2008.
- [3] D. Yuan, Y. Luo, X. Zhuang, G. R. Rodrigues, X. Zhao, Y. Zhang, P. U. Jain, M. Stumm, Simple testing can prevent most critical failures: An analysis of production failures in distributed data-intensive systems, in: 11th USENIX Symposium on Operating Systems Design and Implementation (OSDI), 2014.
- [4] A. Alquraan, H. Takruri, M. Alfatafta, S. Al-Kiswany, An analysis of network-partitioning failures in cloud systems, in: 13th USENIX Symposium on Operating Systems Design and Implementation (OSDI), 2018, pp. 51–68.

- [5] J. Xue, Y. Zhao, Y. Wang, J. Chen, D. Wang, A characterization study of bugs in LLM agent workflow orchestration frameworks, in: 40th IEEE/ACM International Conference on Automated Software Engineering (ASE), Industry Showcase, 2025, pp. 3369–3380. doi:10.1109/ASE63991.2025.00278.
- [6] M. Liu, S. Zhong, W. Bi, Y. Zhang, Z. Chen, Z. Chen, X. Liu, Y. Ma, A first look at bugs in LLM inference engines, *ACM Transactions on Software Engineering and Methodology* (2026). doi:10.1145/3788873.
- [7] W. Yu, et al., Maltracker: A fine-grained NPM malware tracker copiloted by LLM-enhanced dataset, in: Proceedings of the 33rd ACM SIGSOFT International Symposium on Software Testing and Analysis (ISSTA), 2024. doi:10.1145/3650212.3680397.
- [8] W. Zhao, et al., AppPoet: LLM-based Android malware detection via multi-view prompt engineering, *Expert Systems with Applications* 262, arXiv:2404.18816 (2025). doi:10.1016/j.eswa.2024.125546.
- [9] A. Rollinson, N. Polatidis, LLM-generated samples for Android malware detection, *Digital* 6 (1) (2026) 5. doi:10.3390/digital6010005.
- [10] S. Saha, et al., MaLAWare: Automating the comprehension of malicious software behaviours using LLMs, in: IEEE/ACM International Conference on Mining Software Repositories (MSR), 2025, arXiv:2504.01145. doi:10.1109/msr66628.2025.00037.
- [11] Center for Threat-Informed Defense, TRAM: Threat report ATT&CK mapper, Open-source software, MITRE Engenuity, not a paper and cited as the artefact it is; the repository names no preferred citation (2023).
- [12] V. Legoy, M. Caselli, C. Seifert, A. Peter, Automated retrieval of ATT&CK tactics and techniques for cyber threat reports, in: 1st Cyber Threat Intelligence Symposium (CTI), 2020, arXiv:2004.14322. Recorded as a conference contribution by the University of Twente research portal; not indexed under this title in dblp.
- [13] G. Husari, E. Al-Shaer, M. Ahmed, B. Chu, X. Niu, TTPDrill: Automatic and accurate extraction of threat actions from unstructured text of CTI sources, in: Proceedings of the 33rd Annual Computer Security Applications Conference (ACSAC), 2017, pp. 103–115. doi:10.1145/3134600.3134646.
- [14] K. Satvat, R. Gjomemo, V. N. Venkatakrishnan, EXTRACTOR: Extracting attack behavior from threat reports, in: IEEE European Symposium on Security and Privacy (EuroS&P), 2021, pp. 598–615. doi:10.1109/eurosp51992.2021.00046.

- [15] Z. Li, J. Zeng, Y. Chen, Z. Liang, AttacKG: Constructing technique knowledge graph from cyber threat intelligence reports, in: *Computer Security, ESORICS*, 2022, pp. 589–609. doi:10.1007/978-3-031-17140-6_29.
- [16] M. T. Alam, D. Bhusal, Y. Park, N. Rastogi, Looking beyond IoCs: Automatically extracting attack patterns from external CTI, in: *26th International Symposium on Research in Attacks, Intrusions and Defenses (RAID)*, 2023, pp. 92–108. doi:10.1145/3607199.3607208.
- [17] A. Lekssays, U. Shukla, H. T. Sencar, M. Parvez, TechniqueRAG: Retrieval augmented generation for adversarial technique annotation in CTI text, in: *Findings of the Association for Computational Linguistics (ACL)*, 2025, arXiv:2505.11988. doi:10.18653/v1/2025.findings-acl.1076.
- [18] F. Morbiato, M. Keller, P. Nair, L. Romano, Hierarchical retrieval augmented generation for adversarial technique annotation in cyber threat intelligence text, PREPRINT, not peer reviewed. arXiv:2604.14166 (2026). doi:10.48550/arXiv.2604.14166.
- [19] Y. Cheng, C. Li, R. S. P. Basuki, Q. Cui, W. Ding, P. Gao, TTPrint: Evidence-grounded TTP extraction via diverge-then-converge verification, PREPRINT, not peer reviewed. arXiv:2605.25836 (2026). doi:10.48550/arXiv.2605.25836.
- [20] K. D’Oosterlinck, O. Khattab, F. Remy, T. Demeester, C. Develder, C. Potts, In-context learning for extreme multi-label classification, PREPRINT, not peer reviewed. arXiv:2401.12178 (2024). doi:10.48550/arXiv.2401.12178.
- [21] A. Abdallah, J. Holdcroft, M. Ali, A. Jatowt, Are LLM-based retrievers worth their cost? an empirical study of efficiency, robustness, and reasoning overhead, in: *Proceedings of the 49th International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2026, arXiv:2604.03676. doi:10.1145/3805712.3808566.
- [22] L. Lange, H. Adel, et al., AnnoCTR: A dataset for detecting and linking entities, tactics, and techniques in cyber threat reports, in: *Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING)*, 2024, arXiv:2404.07765. CC-BY-SA 4.0. doi:10.63317/4esp6coivous.
- [23] H. T. Ng, et al., The CoNLL-2014 shared task on grammatical error correction, in: *Proceedings of the Eighteenth Conference on Computational Natural Language Learning: Shared Task*, 2014. doi:10.3115/v1/w14-1701.
- [24] C. Bryant, M. Felice, Ø. E. Andersen, T. Briscoe, The BEA-2019 shared task on grammatical error correction, in: *Proceedings of the Fourteenth Workshop*

- on Innovative Use of NLP for Building Educational Applications, 2019. doi: 10.18653/v1/w19-4406.
- [25] S. Welleck, I. Kulikov, S. Roller, E. Dinan, K. Cho, J. Weston, Neural text generation with unlikelihood training, in: International Conference on Learning Representations (ICLR), 2020, arXiv:1908.04319.
 - [26] H. Li, T. Lan, Z. Fu, D. Cai, L. Liu, N. Collier, T. Watanabe, Y. Su, Repetition in repetition out: Towards understanding neural text degeneration from the data perspective, in: Advances in Neural Information Processing Systems (NeurIPS), 2023, arXiv:2310.10226. doi:10.52202/075280-3187.
 - [27] G. Chen, H. Sun, D. Liu, Z. Wang, Q. Wang, B. Yin, L. Liu, L. Ying, Re-Copilot: Reverse engineering copilot in binary analysis, PREPRINT, not peer reviewed. arXiv:2505.16366 (2025). doi:10.48550/arXiv.2505.16366.
 - [28] H. Tan, Q. Luo, J. Li, Y. Zhang, LLM4Decompile: Decompiling binary code with large language models, in: Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP), 2024, pp. 3473–3487. doi:10.18653/v1/2024.emnlp-main.203.
 - [29] N. Jiang, C. Wang, K. Liu, X. Xu, L. Tan, X. Zhang, P. Babkin, Nova: Generative language models for assembly code with hierarchical attention and contrastive learning, in: International Conference on Learning Representations (ICLR), 2025, poster. arXiv:2311.13721.
 - [30] Z. Xuan, X. Xu, M. Zheng, L. Z.-H. Tan, J. Guo, T. Zhang, L. Yu, C. Wang, X. Zhang, Identifying adversary tactics and techniques in malware binaries with an LLM agent, PREPRINT, not peer reviewed. arXiv:2602.06325 (2026). doi:10.48550/arXiv.2602.06325.
 - [31] K. Kate, T. Pedapati, K. Basu, Y. Rizk, V. Chenthamarakshan, S. Chaudhury, M. Agarwal, I. Abdelaziz, LongFuncEval: Measuring the effectiveness of long context models for function calling, PREPRINT, not peer reviewed. arXiv:2505.10570 (2025). doi:10.48550/arXiv.2505.10570.
 - [32] V. Paramanayakam, A. Karatzas, I. Anagnostopoulos, D. Stamoulis, Less is more: Optimizing function calling for LLM execution on edge devices, in: Design, Automation and Test in Europe Conference (DATE), 2025, arXiv:2411.15399. doi:10.23919/date64628.2025.10992798.
 - [33] D. Lea, J. Ghawaly, G. G. Richard III, A. Ali-Gombe, A. Case, REx86: A local large language model for assisting in x86 assembly reverse engineering, in: Annual Computer Security Applications Conference (ACSAC), 2025, arXiv:2510.20975. doi:10.1109/acsac67867.2025.00024.

- [34] W. Li, S. Manickam, Y.-w. Chong, S. Karuppayah, PhishDebate: An LLM-based multi-agent framework for phishing website detection, in: IEEE International Conference on Big Data (BigData), 2025, pp. 6606–6615, arXiv:2506.15656. doi:10.1109/BigData66926.2025.11401440.
- [35] D. Tran, D. Kiela, Single-agent LLMs outperform multi-agent systems on multi-hop reasoning under equal thinking token budgets, PREPRINT, not peer reviewed. arXiv:2604.02460 (2026). doi:10.48550/arXiv.2604.02460.
- [36] B. Bertalaníč, C. Fortuna, The cost of consensus: Isolated self-correction prevails over unguided homogeneous multi-agent debate, in: ACM Conference on AI and Agentic Systems, 2026, arXiv:2605.00914. doi:10.1145/3786335.3813137.
- [37] R. Evertz, N. Risse, C. Neuer, N. Müller, S. Normann, et al., Chasing shadows: Pitfalls in LLM security research, in: Network and Distributed System Security Symposium (NDSS), 2026, arXiv:2512.09549. doi:10.14722/ndss.2026.241749.
- [38] J. Ng, A. Milani Fard, Evaluating retrieval-augmented generation for explainable malware analysis, in: ACM SecDev, 2026, poster. arXiv:2605.03140.
- [39] J. Metz, L. F. D. de Abreu, E. A. Cherman, M. C. Monard, On the estimation of predictive evaluation measure baselines for multi-label learning, in: Advances in Artificial Intelligence, IBERAMIA, Lecture Notes in Computer Science, 2012, pp. 189–198. doi:10.1007/978-3-642-34654-5_20.
- [40] J. Metz, N. Spolaôr, E. A. Cherman, M. C. Monard, Comparing published multi-label classifier performance measures to the ones obtained by a simple multi-label baseline classifier, PREPRINT, not peer reviewed. arXiv:1503.06952 (2015). doi:10.48550/arXiv.1503.06952.
- [41] A. Lazaridou, A. Kuncoro, E. Gribovskaya, et al., Mind the gap: Assessing temporal generalization in neural language models, in: Advances in Neural Information Processing Systems (NeurIPS), 2021, arXiv:2102.01951.
- [42] T. Xu, J. Zhang, P. Huang, J. Zheng, T. Sheng, D. Yuan, Y. Zhou, S. Pasupathy, Do not blame users for misconfigurations, in: 24th ACM Symposium on Operating Systems Principles (SOSP), 2013, pp. 244–259. doi:10.1145/2517349.2522727.
- [43] O. Oyelayo, S. Abedu, S. Khatoonabadi, E. Shihab, Developer challenges in the adoption of model context protocol in AI agent development, in: 7th International Workshop on Bots and Agents in Software Engineering (BotSE), co-located with ICSE, 2026, pp. 53–59. doi:10.1145/3786161.3788462.

- [44] A. Martin-Lopez, S. Segura, A. Ruiz-Cortés, A catalogue of inter-parameter dependencies in RESTful web APIs, in: *Service-Oriented Computing (IC-SOC)*, 2019, pp. 399–414. doi:10.1007/978-3-030-33702-5_31.
- [45] Y. Li, Z. Liu, P.-L. Poon, D. Towey, C.-A. Sun, et al., Metamorphic relation generation: State of the art and research directions, *ACM Transactions on Software Engineering and Methodology* ArXiv:2406.05397 (2024). doi:10.1145/3708521.
- [46] S. Dev, A. Sloan, J. Kavner, N. Kong, M. Sandler, Judge reliability harness: Stress testing the reliability of LLM judges, in: *Agents in the Wild: Safety, Security, and Beyond Workshop, International Conference on Learning Representations (ICLR)*, 2026, arXiv:2603.05399.
- [47] A. Ryan, S. S. Noor, M. Erfan, S. Mitra, S. Mittal, M. R. Rahman, Evaluating open-source LLMs for multi-label ATT&CK technique classification on CTI reports, PREPRINT, not peer reviewed. arXiv:2606.18166 (2026). doi:10.48550/arXiv.2606.18166.
- [48] Y. Sun, H. Wang, J. Li, J. Liu, X. Li, H. Wen, Y. Yuan, H. Zheng, Y. Liang, Y. Li, Y. Liu, An empirical study of LLM reasoning ability under strict output length constraint, in: *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, 2025, pp. 7652–7671. doi:10.18653/v1/2025.emnlp-main.389.
- [49] M. Mizrahi, G. Kaplan, D. Malkin, R. Dror, D. Shahaf, G. Stanovsky, State of what art? A call for multi-prompt LLM evaluation, *Transactions of the Association for Computational Linguistics* 12 (2024) 933–949. doi:10.1162/tac1_a_00681.
- [50] C. Snell, J. Lee, K. Xu, A. Kumar, Scaling LLM test-time compute optimally can be more effective than scaling parameters for reasoning, in: *International Conference on Learning Representations (ICLR)*, 2025, oral.
- [51] M. Hassid, T. Remez, J. Gehring, R. Schwartz, Y. Adi, The larger the better? Improved LLM code-generation via budget reallocation, in: *First Conference on Language Modeling (COLM)*, 2024.
- [52] M. Jin, Q. Yu, D. Shu, H. Zhao, W. Hua, Y. Meng, Y. Zhang, M. Du, The impact of reasoning step length on large language models, in: *Findings of the Association for Computational Linguistics: ACL*, 2024, pp. 1830–1842. doi:10.18653/v1/2024.findings-acl.108.
- [53] D. Küçük, F. Ertam, Replication package: the pipeline, the evaluation harnesses, the per-sample records and the derivation of every number reported, Source repository, under an open licence. <https://github.com/Root0ne/Maljan> (2026).

- [54] D. Kumaran, Reported confidence in LLMs tracks commitment more than correctness, PREPRINT, not peer reviewed. arXiv:2606.29490 (2026). doi: 10.48550/arXiv.2606.29490.
- [55] ggml-org/llama.cpp, max_tokens not respected on the non-chat endpoint, PREPRINT, not peer reviewed. ggml-org/llama.cpp issue #8634 (2024).
- [56] abuse.ch, MalwareBazaar, Malware sample corpus. <https://bazaar.abuse.ch>, cited as the corpus the dynamic cohort was drawn from. Its terms permit research redistribution of individual samples but not of the corpus as a package (2026).
- [57] The MITRE Corporation, MITRE ATT&CK enterprise, Knowledge base. <https://attack.mitre.org>, the family-to-technique uses relationships used here as per-sample ground truth are derived from the mitre-attack/attack-stix-data distribution, under the ATT&CK Terms of Use (2026).